



Multi-Source Market Intelligence System:

This notebook builds an end-to-end data analysis pipeline using:

- Market price data (API)
- News headlines (GDELT)
- Sentiment-based feature engineering
- Time-series analysis and insights

In [300...]: `!pip install pandas numpy requests matplotlib seaborn vaderSentiment`

```
Requirement already satisfied: pandas in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (2.3.3)
Requirement already satisfied: numpy in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (2.3.2)
Requirement already satisfied: requests in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (2.32.5)
Requirement already satisfied: matplotlib in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (3.10.6)
Requirement already satisfied: seaborn in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (0.13.2)
Requirement already satisfied: vaderSentiment in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (3.3.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from pandas) (2025.2)
Requirement already satisfied: charset_normalizer<4,>=2 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from requests) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from requests) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from requests) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from requests) (2025.10.5)
Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from matplotlib) (4.59.2)
Requirement already satisfied: kiwisolver>=1.3.1 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from matplotlib) (3.2.3)
Requirement already satisfied: six>=1.5 in /opt/anaconda3/envs/Assignments/lib/python3.13/site-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
```

In [301...]

```
import requests #lets python talk to the internet
import pandas as pd
import numpy as np

import urllib.parse

import matplotlib.pyplot as plt
import seaborn as sns

from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer #Import rule-based sentiment analyzer: doesn't train model but uses predefined
```

◆ STEP 1: ENVIRONMENT SETUP & LIBRARY IMPORTS

```
In [302... pd.set_option("Display.max_columns", None) #Show all columns while displaying
pd.set_option("Display.float_format","{:.4f}".format) #Format floating numbers

plt.style.use("seaborn-v0_8") #Pre-defined plotting theme
sns.set_context("talk") #Scales plot elements (paper = small, notebook = medium)

print("Environment Setup Complete!")
```

Environment Setup Complete!

◆ STEP 2: FETCH MARKET DATA (CoinGecko API)

In this step, we fetch historical cryptocurrency market data. This data will serve as the backbone for our time-series analysis.

```
In [303... url = "https://api.coingecko.com/api/v3/coins/bitcoin/market_chart"

params = {
    "vs_currency": "usd", #Prices in USD
    "days": 180, #Last 180 days
    "interval": "daily" #daily data
}

response = requests.get(url, params = params) #API request

response.status_code #Anything other than 200 stop and debug
```

Out[303... 200

```
In [304... data = response.json() #Formats our response to JSON

data.keys() #Return keys of our data
```

Out[304... dict_keys(['prices', 'market_caps', 'total_volumes'])

```
In [305... prices = data["prices"]
print(prices[:5]) #[timestamp, value] -> time and bitcoin price at that time

[[1752364800000, 117418.95745007684], [1752451200000, 119117.55666327637],
[1752537600000, 119833.67446712355], [1752624000000, 117678.19493404306],
[1752710400000, 118748.1627367753]]
```

```
In [306... price_df = pd.DataFrame(prices, columns=["timestamp","price"]) #converts
price_df.head()
```

	timestamp	price
0	1752364800000	117418.9575
1	1752451200000	119117.5567
2	1752537600000	119833.6745
3	1752624000000	117678.1949
4	1752710400000	118748.1627

```
In [307...]: #Convert timestamp to date
price_df["date"] = pd.to_datetime(price_df["timestamp"], unit="ms").dt.date
price_df = price_df[["date","price"]]
price_df.head()
```

```
Out[307...]:
```

	date	price
0	2025-07-13	117418.9575
1	2025-07-14	119117.5567
2	2025-07-15	119833.6745
3	2025-07-16	117678.1949
4	2025-07-17	118748.1627

```
In [308...]: price_df.shape
price_df.info()
price_df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 181 entries, 0 to 180
Data columns (total 2 columns):
 #   Column   Non-Null Count   Dtype  
 ---  --       --           --      
 0   date     181 non-null    object 
 1   price    181 non-null    float64
dtypes: float64(1), object(1)
memory usage: 3.0+ KB
```

```
Out[308...]:
```

	price
count	181.0000
mean	106233.7453
std	12039.7923
min	84682.6243
25%	92036.7255
50%	110708.6696
75%	115970.5849
max	124773.5082

```
In [309...]: price_df.isnull().sum()
```

```
Out[309...]:
```

date	0
price	0
	dtype: int64

```
In [310...]: price_df.duplicated().sum()
```

```
Out[310...]: np.int64(0)
```

```
In [311...]: price_df["date"].min(), price_df["date"].max()
```

```
Out[311... (datetime.date(2025, 7, 13), datetime.date(2026, 1, 8))
```

```
In [312... price_df["price"] = price_df["price"].interpolate()  
price_df = price_df.drop_duplicates()  
price_df.isnull().sum()
```

```
Out[312... date      0  
price     0  
dtype: int64
```

◆ STEP 3: NEWS DATA COLLECTION (GDELT)

```
In [313... news_url = "https://api.gdeltproject.org/api/v2/doc/doc"  
  
query = "(bitcoin OR cryptocurrency OR crypto market)" #News articles may  
start_date = "20250701000000"    # YYYYMMDDHHMMSS  
end_date   = "20260101000000"  
  
news_params = {  
    "query": query,  
    "mode": "ArtList", #Return list of articles  
    "format": "JSON",  
    "maxrecords": 250, #Max articles per call  
    "startdate": start_date,  
    "enddate": end_date  
}
```

```
In [314... url = news_url + "?" + urllib.parse.urlencode(news_params) #to see what w  
url
```

```
Out[314... 'https://api.gdeltproject.org/api/v2/doc/doc?query=%28bitcoin+OR+cryptoc  
urrency+OR+crypto+market%29&mode=ArtList&format=JSON&maxrecords=250&star  
tdatetime=20250701000000&enddatetime=20260101000000'
```

```
In [315... news_response = requests.get(url)  
  
news_response.status_code
```

```
Out[315... 200
```

```
In [316... news_data = news_response.json()  
  
news_data.keys()
```

```
Out[316... dict_keys(['articles'])
```

```
In [317... articles = news_data["articles"]  
  
len(articles) #how many news articles we fetched in this API
```

```
Out[317... 250
```

```
In [318... news_df = pd.DataFrame(articles)  
  
news_df.head()
```

```
Out[318...]
```

url

0 https://finance.yahoo.com/news/buy-bitcoin-whi...

1 https://www.theguardian.com/technology/2025/de...

2 https://finance.yahoo.com/news/bitcoin-ethereu...

3 https://charter97.org/en/news/2025/11/6/661980/ https://qnt91x.c97.org/en/news/2025

4 https://finance.yahoo.com/news/2-best-cryptocu...

```
In [319...]: news_df = news_df[["seendate", "title", "domain"]]
```

```
In [320...]: news_df["date"] = pd.to_datetime(news_df["seendate"]).dt.date  
news_df = news_df.drop(columns=["seendate"])
```

```
news_df.head()
```

```
Out[320...]
```

		title	domain	date
0	Should You Buy Bitcoin While It Under \$100 , 0...	finance.yahoo.com	2025-12-07	
1	Cryptocurrency slump erases 2025 financial gai...	theguardian.com	2025-12-29	
2	Bitcoin , Ethereum , Dogecoin Trade Flat , Whi...	finance.yahoo.com	2025-12-24	
3	Trump : Were Turning The U . S . Into A Bitcoi...	charter97.org	2025-11-06	
4	The 2 Best Cryptocurrencies to Buy With \$100 R...	finance.yahoo.com	2025-12-20	

```
In [321...]:
```

```
news_df.shape  
news_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 250 entries, 0 to 249  
Data columns (total 3 columns):  
 #   Column  Non-Null Count  Dtype    
---  --    
 0   title    250 non-null    object  
 1   domain   250 non-null    object  
 2   date     250 non-null    object  
dtypes: object(3)  
memory usage: 6.0+ KB
```

```
In [322...]:
```

```
news_df.isnull().sum()
```

```
Out[322... title      0
domain      0
date        0
dtype: int64

In [323... (news_df["title"].str.strip() == "") .sum() #check for empty or blank titl
Out[323... np.int64(0)

In [324... news_df.duplicated(subset = ["title","date"]).sum()

Out[324... np.int64(29)

In [325... news_df["date"].min(),news_df["date"].max()

Out[325... (datetime.date(2025, 10, 3), datetime.date(2026, 1, 1))

In [326... news_df = news_df.reset_index(drop = True) #reset the index of DataFrame
```

◆ STEP 4: SENTIMENT FEATURE ENGINEERING (VADER)

In this step, we convert news headlines into numeric sentiment scores using a rule-based sentiment analyzer. These scores will later be aggregated and aligned with market data.

```
In [327... analyzer = SentimentIntensityAnalyzer()

In [328... sample_text = news_df.loc[0,"title"]
sample_text

Out[328... 'Should You Buy Bitcoin While It Under $100 , 000 ? '

In [329... analyzer.polarity_scores(sample_text)
# neg = negative, neu = neutral, pos = positive, compound = overall senti
#Compound: >= +0.05 positive
#           <= -0.05 negative
#           between neutral

Out[329... {'neg': 0.0, 'neu': 1.0, 'pos': 0.0, 'compound': 0.0}

In [330... news_df["sentiment_dict"] = news_df["title"].apply(
    lambda x: analyzer.polarity_scores(x)
)

In [331... news_df.head()
```

Out[331...]

	title	domain	date	sentiment_dict
0	Should You Buy Bitcoin While It Under \$100 , 0...	finance.yahoo.com	2025-12-07	{'neg': 0.0, 'neu': 1.0, 'pos': 0.0, 'compound': 0.468}
1	Cryptocurrency slump erases 2025 financial gai...	theguardian.com	2025-12-29	{'neg': 0.0, 'neu': 0.468, 'pos': 0.532, 'compound': 0.532}
2	Bitcoin , Ethereum , Dogecoin Trade Flat , Whi...	finance.yahoo.com	2025-12-24	{'neg': 0.0, 'neu': 0.891, 'pos': 0.109, 'compound': 0.891}
3	Trump : Were Turning The U . S . Into A Bitcoi...	charter97.org	2025-11-06	{'neg': 0.0, 'neu': 1.0, 'pos': 0.0, 'compound': 0.468}
4	The 2 Best Cryptocurrencies to Buy With \$100 R...	finance.yahoo.com	2025-12-20	{'neg': 0.0, 'neu': 0.682, 'pos': 0.318, 'compound': 0.682}

In [332...]

```
news_df["sentiment_neg"] = news_df["sentiment_dict"].apply(lambda x: x["neg"])
news_df["sentiment_neu"] = news_df["sentiment_dict"].apply(lambda x: x["neu"])
news_df["sentiment_pos"] = news_df["sentiment_dict"].apply(lambda x: x["pos"])
news_df["sentiment_compound"] = news_df["sentiment_dict"].apply(lambda x: x["compound"])

news_df = news_df.drop(columns=["sentiment_dict"])
```

In [333...]

news_df.head()

Out[333...]

	title	domain	date	sentiment_neg	sentiment_neu	sentiment_pos	sentiment_compound
0	Should You Buy Bitcoin While It Under \$100 , 0...	finance.yahoo.com	2025-12-07	0.0000	1.0000	0.0000	0.4680
1	Cryptocurrency slump erases 2025 financial gai...	theguardian.com	2025-12-29	0.0000	0.4680	0.5320	0.5320
2	Bitcoin , Ethereum , Dogecoin Trade Flat , Whi...	finance.yahoo.com	2025-12-24	0.0000	0.8910	0.1090	0.8910
3	Trump : Were Turning The U . S . Into A Bitcoi...	charter97.org	2025-11-06	0.0000	1.0000	0.0000	0.4680
4	The 2 Best Cryptocurrencies to Buy With \$100 R...	finance.yahoo.com	2025-12-20	0.0000	0.6820	0.3180	0.6820

◆ STEP 5: DAILY SENTIMENT AGGREGATION

As our market price data has price per day whereas news consists of various articles every day so we need to aggregate the articles per day in our news DataFrame

In [334...]

news_df["date"].value_counts().head()

```
Out[334... date
2025-12-01    21
2025-11-24    11
2025-12-02     9
2025-12-29     8
2025-11-01     8
Name: count, dtype: int64
```

```
In [335... daily_sentiment = (
    news_df.groupby("date").agg(
        avg_compound=("sentiment_compound", "mean"),
        avg_positive=("sentiment_pos", "mean"),
        avg_negative=("sentiment_neg", "mean"),
        avg_neutral=("sentiment_neu", "mean"),
        article_count=("sentiment_compound", "count") #News volume
    )
    .reset_index() #Converts date to back to column
)

daily_sentiment.head()
```

```
Out[335...   date  avg_compound  avg_positive  avg_negative  avg_neutral  article_count
0  2025-10-03      0.1327      0.1123      0.0313      0.8564          7
1  2025-10-04      0.2533      0.1202      0.0000      0.8798          4
2  2025-10-05     -0.0512      0.0000      0.0454      0.9546          5
3  2025-10-06      0.1346      0.1295      0.0580      0.8125          2
4  2025-10-08      0.1027      0.1360      0.1170      0.7480          1
```

◆ STEP 5: MERGING MARKET DATA WITH DAILY SENTIMENT

In this step, we merge daily market price data with aggregated daily news sentiment to create a unified analysis-ready dataset.

```
In [336... price_df.head()
```

```
Out[336...   date      price
0  2025-07-13  117418.9575
1  2025-07-14  119117.5567
2  2025-07-15  119833.6745
3  2025-07-16  117678.1949
4  2025-07-17  118748.1627
```

```
In [337...]: daily_sentiment.head()
```

```
Out[337...]:
```

	date	avg_compound	avg_positive	avg_negative	avg_neutral	article_count
0	2025-10-03	0.1327	0.1123	0.0313	0.8564	7
1	2025-10-04	0.2533	0.1202	0.0000	0.8798	4
2	2025-10-05	-0.0512	0.0000	0.0454	0.9546	5
3	2025-10-06	0.1346	0.1295	0.0580	0.8125	2
4	2025-10-08	0.1027	0.1360	0.1170	0.7480	1

```
In [338...]: #Merging on date column
```

```
final_df = price_df.merge(daily_sentiment, on = "date", how = "left")
```

```
# Why left? : keep all market date and attach sentiment only if available
```

```
final_df.head()
```

```
Out[338...]:
```

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral	a
0	2025-07-13	117418.9575	NaN	NaN	NaN	NaN	NaN
1	2025-07-14	119117.5567	NaN	NaN	NaN	NaN	NaN
2	2025-07-15	119833.6745	NaN	NaN	NaN	NaN	NaN
3	2025-07-16	117678.1949	NaN	NaN	NaN	NaN	NaN
4	2025-07-17	118748.1627	NaN	NaN	NaN	NaN	NaN

```
In [339...]: final_df.tail()
```

```
Out[339...]:
```

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral
176	2026-01-05	91373.2159	NaN	NaN	NaN	NaN
177	2026-01-06	93926.7956	NaN	NaN	NaN	NaN
178	2026-01-07	93666.8639	NaN	NaN	NaN	NaN
179	2026-01-08	91257.1554	NaN	NaN	NaN	NaN
180	2026-01-08	89788.8330	NaN	NaN	NaN	NaN

```
In [340...]: final_df.isnull().sum() #sentiment are null as there is no relevant news
```

```
Out[340...]: date          0  
price          0  
avg_compound  108  
avg_positive  108  
avg_negative  108  
avg_neutral   108  
article_count  108  
dtype: int64
```

```
In [341...]: final_df.info()  
final_df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 181 entries, 0 to 180  
Data columns (total 7 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   date             181 non-null    object    
 1   price            181 non-null    float64  
 2   avg_compound    73 non-null    float64  
 3   avg_positive    73 non-null    float64  
 4   avg_negative    73 non-null    float64  
 5   avg_neutral     73 non-null    float64  
 6   article_count   73 non-null    float64  
dtypes: float64(6), object(1)  
memory usage: 10.0+ KB
```

```
Out[341...]:      price  avg_compound  avg_positive  avg_negative  avg_neutral  artic  
count      181.0000      73.0000      73.0000      73.0000      73.0000  
mean     106233.7453      0.0422      0.0718      0.0610      0.8672  
std      12039.7923      0.2467      0.0740      0.0671      0.0792  
min      84682.6243     -0.6249      0.0000      0.0000      0.6985  
25%     92036.7255     -0.0705      0.0000      0.0000      0.8160  
50%     110708.6696      0.0169      0.0557      0.0510      0.8690  
75%     115970.5849      0.2001      0.1117      0.0920      0.9187  
max     124773.5082      0.5973      0.3015      0.2856      1.0000
```

◆ STEP 6: HANDLING MISSING SENTIMENT VALUES

After merging market and sentiment data, some days contain missing sentiment values due to the absence of news articles. These are handled explicitly to preserve semantic meaning.

```
In [342...]: final_df[final_df["avg_compound"].isnull()].head()
```

Out[342...]

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral	a
0	2025-07-13	117418.9575	NaN	NaN	NaN	NaN	NaN
1	2025-07-14	119117.5567	NaN	NaN	NaN	NaN	NaN
2	2025-07-15	119833.6745	NaN	NaN	NaN	NaN	NaN
3	2025-07-16	117678.1949	NaN	NaN	NaN	NaN	NaN
4	2025-07-17	118748.1627	NaN	NaN	NaN	NaN	NaN

In [343...]

```
#Instead of dropping missing values from sentiment, we can replace null with 0
final_df["avg_compound"] = final_df["avg_compound"].fillna(0)
final_df["avg_positive"] = final_df["avg_positive"].fillna(0)
final_df["avg_negative"] = final_df["avg_negative"].fillna(0)
final_df["avg_neutral"] = final_df["avg_neutral"].fillna(0)
final_df["article_count"] = final_df["article_count"].fillna(0)
```

In [344...]

```
final_df.isnull().sum()
```

Out[344...]

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral	article_count
	0	0	0	0	0	0	0
	1	0	0	0	0	0	0
	2	0	0	0	0	0	0
	3	0	0	0	0	0	0
	4	0	0	0	0	0	0

In [345...]

```
final_df.head()
```

Out[345...]

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral	a
0	2025-07-13	117418.9575	0.0000	0.0000	0.0000	0.0000	0.0000
1	2025-07-14	119117.5567	0.0000	0.0000	0.0000	0.0000	0.0000
2	2025-07-15	119833.6745	0.0000	0.0000	0.0000	0.0000	0.0000
3	2025-07-16	117678.1949	0.0000	0.0000	0.0000	0.0000	0.0000
4	2025-07-17	118748.1627	0.0000	0.0000	0.0000	0.0000	0.0000

In [346...]

```
final_df.tail()
```

Out [346...]

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral
176	2026-01-05	91373.2159	0.0000	0.0000	0.0000	0.0000
177	2026-01-06	93926.7956	0.0000	0.0000	0.0000	0.0000
178	2026-01-07	93666.8639	0.0000	0.0000	0.0000	0.0000
179	2026-01-08	91257.1554	0.0000	0.0000	0.0000	0.0000
180	2026-01-08	89788.8330	0.0000	0.0000	0.0000	0.0000

◆ STEP 7: CREATING SENTIMENT REGIMES

In this step, we categorize daily sentiment scores into interpretable sentiment regimes (Positive, Neutral, Negative) to enable regime-based analysis. A regime is a state or condition of the system.

In [347...]

```
# creating a new column to explicitly define type of sentiment on the basis of avg_compound score
def classify_sentiment(score):
    if score >= 0.05:
        return "Positive"
    elif score <= -0.05:
        return "Negative"
    else:
        return "Neutral"
```

In [348...]

```
final_df[["sentiment_regime"]] = final_df[["avg_compound"]].apply(classify_sentiment)
final_df.head()
```

Out [348...]

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral	a
0	2025-07-13	117418.9575	0.0000	0.0000	0.0000	0.0000	0.0000
1	2025-07-14	119117.5567	0.0000	0.0000	0.0000	0.0000	0.0000
2	2025-07-15	119833.6745	0.0000	0.0000	0.0000	0.0000	0.0000
3	2025-07-16	117678.1949	0.0000	0.0000	0.0000	0.0000	0.0000
4	2025-07-17	118748.1627	0.0000	0.0000	0.0000	0.0000	0.0000

In [349...]

```
final_df[["sentiment_regime"]].value_counts()
```

```
Out[349...]: sentiment_regime  
Neutral      127  
Positive      31  
Negative      23  
Name: count, dtype: int64
```

◆ STEP 8: PRICE RETURNS & VOLATILITY FEATURE ENGINEERING

In this step, we transform raw price data into daily returns and volatility measures to enable meaningful market analysis.

- returns → how much price changes
- volatility → how unstable price is

```
In [350...]: final_df = final_df.sort_values("date").reset_index(drop = True)
```

```
In [351...]: final_df["daily_return"] = final_df["price"].pct_change()  
  
# Percentage Change  
# pct = (price[t] - price[t-1])/ price[t-1]  
  
final_df[["date","price","daily_return"]].head()  
# first day return is NaN as there is no previous day
```

```
Out[351...]:
```

	date	price	daily_return
0	2025-07-13	117418.9575	NaN
1	2025-07-14	119117.5567	0.0145
2	2025-07-15	119833.6745	0.0060
3	2025-07-16	117678.1949	-0.0180
4	2025-07-17	118748.1627	0.0091

```
In [352...]: # Create rolling volatility (7-day)  
final_df["volatility_7d"] = (final_df["daily_return"].rolling(window=7).s  
  
# What is Volatility: Markets often react to negative news with volatility  
# Higher volatility → more uncertainty && Lower volatility → stability  
  
final_df["volatility_14d"] = (final_df["daily_return"].rolling(window=14).s  
  
final_df["volatility_30d"] = (final_df["daily_return"].rolling(window=30).s
```

```
In [353...]: final_df.head(31)  
  
# daily_return start at second day  
# volatility start at 8th day with rolling volatility for the first 7 day
```

Out[353...]

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral
0	2025-07-13	117418.9575	0.0000	0.0000	0.0000	0.0000
1	2025-07-14	119117.5567	0.0000	0.0000	0.0000	0.0000
2	2025-07-15	119833.6745	0.0000	0.0000	0.0000	0.0000
3	2025-07-16	117678.1949	0.0000	0.0000	0.0000	0.0000
4	2025-07-17	118748.1627	0.0000	0.0000	0.0000	0.0000
5	2025-07-18	119445.3652	0.0000	0.0000	0.0000	0.0000
6	2025-07-19	117988.9466	0.0000	0.0000	0.0000	0.0000
7	2025-07-20	117901.6266	0.0000	0.0000	0.0000	0.0000
8	2025-07-21	117256.9208	0.0000	0.0000	0.0000	0.0000
9	2025-07-22	117482.4698	0.0000	0.0000	0.0000	0.0000
10	2025-07-23	119955.7957	0.0000	0.0000	0.0000	0.0000
11	2025-07-24	118629.0559	0.0000	0.0000	0.0000	0.0000
12	2025-07-25	118354.4352	0.0000	0.0000	0.0000	0.0000
13	2025-07-26	117540.8084	0.0000	0.0000	0.0000	0.0000
14	2025-07-27	117959.5423	0.0000	0.0000	0.0000	0.0000
15	2025-07-28	119418.9141	0.0000	0.0000	0.0000	0.0000
16	2025-07-29	118003.3020	0.0000	0.0000	0.0000	0.0000
17	2025-07-30	117853.3089	0.0000	0.0000	0.0000	0.0000
18	2025-07-31	117833.2409	0.0000	0.0000	0.0000	0.0000
19	2025-08-01	115700.0024	0.0000	0.0000	0.0000	0.0000
20	2025-08-02	113234.6051	0.0000	0.0000	0.0000	0.0000

	date	price	avg_compound	avg_positive	avg_negative	avg_neutral
21	2025-08-03	112554.9023	0.0000	0.0000	0.0000	0.0000
22	2025-08-04	114199.1097	0.0000	0.0000	0.0000	0.0000
23	2025-08-05	115138.6861	0.0000	0.0000	0.0000	0.0000
24	2025-08-06	114128.3541	0.0000	0.0000	0.0000	0.0000
25	2025-08-07	115022.0958	0.0000	0.0000	0.0000	0.0000
26	2025-08-08	117463.4745	0.0000	0.0000	0.0000	0.0000
27	2025-08-09	116688.3666	0.0000	0.0000	0.0000	0.0000
28	2025-08-10	116510.0839	0.0000	0.0000	0.0000	0.0000
29	2025-08-11	119266.9252	0.0000	0.0000	0.0000	0.0000
30	2025-08-12	118773.7996	0.0000	0.0000	0.0000	0.0000

◆ STEP 9: EXPLORATORY DATA ANALYSIS (EDA)

In [354...]


```
# A. Price Trend over Time
plt.figure(figsize = (12,5))

plt.plot(final_df[ "date" ], final_df[ "price" ])

plt.title("Bitcoin Price Over Time")
plt.xlabel("Date")
plt.ylabel("Price ($)")

plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



Analysis:

- The price shows clear periods of upward and downward movement.
- The market is non-stationary (it trends over time).
- This confirms why analyzing returns instead of raw prices is necessary.

In [355]:


```
# B. Daily Returns distribution

plt.figure(figsize = (8,5))

sns.histplot(final_df["daily_return"].dropna(), bins=40, kde=True)

plt.title("Distribution of daily Returns")
plt.xlabel("Daily Return")
plt.ylabel("Frequency")

plt.tight_layout()
plt.show()
```



Analysis:

- Returns are centered around zero (price did not change), Most daily changes are small, A few extreme values exist on both sides.
- The market usually moves moderately but occasionally experiences large shocks.
- This is typical of financial markets and indicates risk concentration in extreme events.


```
In [356...]: # C. Sentiment Over Time
plt.figure(figsize = (12,5))

plt.plot(final_df["date"], final_df["avg_compound"])
plt.axhline(0, linestyle = "--", color ="pink")

plt.title("Average Daily News Sentiment Over Time")
plt.xlabel("Date")
plt.ylabel("Average Sentiment")

plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



Analysis:

- News sentiment is event-driven, not continuous.
- Major events strongly influence sentiment on specific days.

```
In [357...]: # D. Volatility Over Time
plt.figure(figsize=(12,5))

plt.plot(final_df["date"],final_df["volatility_7d"])

plt.title("7-Day Rolling Volatility")
plt.xlabel("Date")
plt.ylabel("Volatility")

plt.tight_layout()
plt.show()
```



Analysis:

- Volatility clusters in certain periods, Calm periods are followed by spikes in volatility.
- Market uncertainty is not constant.
- Risk tends to increase during specific periods rather than randomly, This suggests external factors (like news) may influence instability.


In [358...]

```
# E. Price vs Sentiment
fig, ax1 = plt.subplots(figsize=(12,5))

ax1.plot(final_df["date"], final_df["price"], color = "blue")
ax1.set_xlabel("Date")
ax1.set_ylabel("Price ($)", color= "blue")
ax1.tick_params(axis="y", labelcolor="blue")

ax2 = ax1.twinx()
ax2.plot(final_df["date"], final_df["avg_compound"], color = "red", alpha=0.5)
ax2.set_ylabel("Average Sentiment", color= "red")
ax2.tick_params(axis="y", labelcolor="red")

plt.title("Price and News Sentiment Over Time")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Analysis:

- Some sentiment spikes align with price movements. In other cases, price moves without strong sentiment change.
- News sentiment is not the sole driver of price.
- This supports a descriptive, not causal interpretation.

In [359]:


```
# F. Positive vs Negative Sentiment
plt.figure(figsize=(12,5))

plt.plot(final_df["date"],final_df["avg_positive"], label="Positive", alpha=0.5)
plt.plot(final_df["date"],final_df["avg_negative"], label="Negative", alpha=0.5)

plt.title("Positive vs Negative News Sentiment Over Time")
plt.xlabel("Date")
plt.ylabel("Sentiment Intensity")

plt.legend()
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



Analysis:

- Positive and negative sentiment vary independently.
- Positive sentiment does not always dominate even during price increases.
- Compound sentiment can hide sentiment imbalance.
- Negative tone often appears stronger and more volatile.
- Separating sentiment components provides deeper insight than using a single score.

🧠 Conclusion

News sentiment does not consistently predict returns, but negative sentiment regimes are associated with higher market volatility, indicating increased risk during adverse news cycles.

◆ STEP 10: COMPARING RETURNS & VOLATILITY ACROSS SENTIMENT REGIMES

compare returns and volatility across Positive, Neutral, and Negative sentiment regimes.

```
In [360...]: # A. Average Daily Return by Sentiment Regime
regime_returns = (
    final_df.groupby("sentiment_regime") ["daily_return"].mean()
)

regime_returns

# Average direction of market movement
# Whether negative days tend to have worse returns
# Whether positive days outperform neutral days
```

```
Out[360...]: sentiment_regime
Negative   -0.0017
Neutral    -0.0014
Positive   -0.0007
Name: daily_return, dtype: float64
```

Analysis:

- Average daily returns differ slightly across sentiment regimes.
- Positive sentiment days tend to have marginally higher returns.
- Negative sentiment days often show weaker or negative average returns.
- News sentiment has some association with market direction, but the effect is modest.
- Sentiment alone is not a strong predictor of daily returns.

```
In [361...]: # B. Return Volatility by Sentiment Regime
regime_volatility = (
    final_df.groupby("sentiment_regime") ["daily_return"].std()
)

regime_volatility

# Risk is often more important than return
# Negative news may increase uncertainty
# Volatility reveals stress, not direction
```

```
Out[361...]: sentiment_regime
Negative   0.0230
Neutral    0.0181
Positive   0.0187
Name: daily_return, dtype: float64
```

Analysis:

- Volatility measures how much and how often prices move up and down.
- High volatility → prices change a lot, Low volatility → prices change a little
- Volatility is noticeably higher during negative sentiment regimes.
- Negative news environments are associated with greater market uncertainty.
- Sentiment impacts risk more than return direction.

```
In [362...]: # C. Rolling Volatility Comparison (7-Day)
regime_vol_7d = (
    final_df
    .groupby("sentiment_regime") ["volatility_7d"]
    .mean()
)
regime_vol_7d

# Sustained risk levels across regimes
# Whether bad sentiment leads to prolonged instability
```

```
Out[362...]: sentiment_regime
Negative    0.0201
Neutral     0.0166
Positive    0.0189
Name: volatility_7d, dtype: float64
```

Analysis:

- Rolling volatility remains elevated during negative sentiment periods.
- Negative sentiment has a persistent effect on market risk, not just a one-day impact.

```
In [363...]: regime_summary = pd.DataFrame({
    "avg_daily_return": regime_returns,
    "return_volatility": regime_volatility,
    "avg_7d_volatility": regime_vol_7d
})

regime_summary
```

	avg_daily_return	return_volatility	avg_7d_volatility
sentiment_regime			
Negative	-0.0017	0.0230	0.0201
Neutral	-0.0014	0.0181	0.0166
Positive	-0.0007	0.0187	0.0189


Analysis:

- Negative sentiment regimes combine lower returns with higher volatility.
- Positive regimes show slightly better stability.
- Market behavior clearly differs by sentiment regime.
- Regime-based analysis provides clearer insight than looking at averages alone.

```
In [364...]: regime_summary.plot(kind = "bar", figsize = (12,5))

plt.title("Market Behavior Across Sentiment Regimes")
plt.xlabel("Value")

plt.tight_layout()
plt.show()
```



🧠 Conclusion

Negative sentiment regimes are associated with higher market volatility, while returns vary only modestly across sentiment states, suggesting sentiment affects risk more than direction.